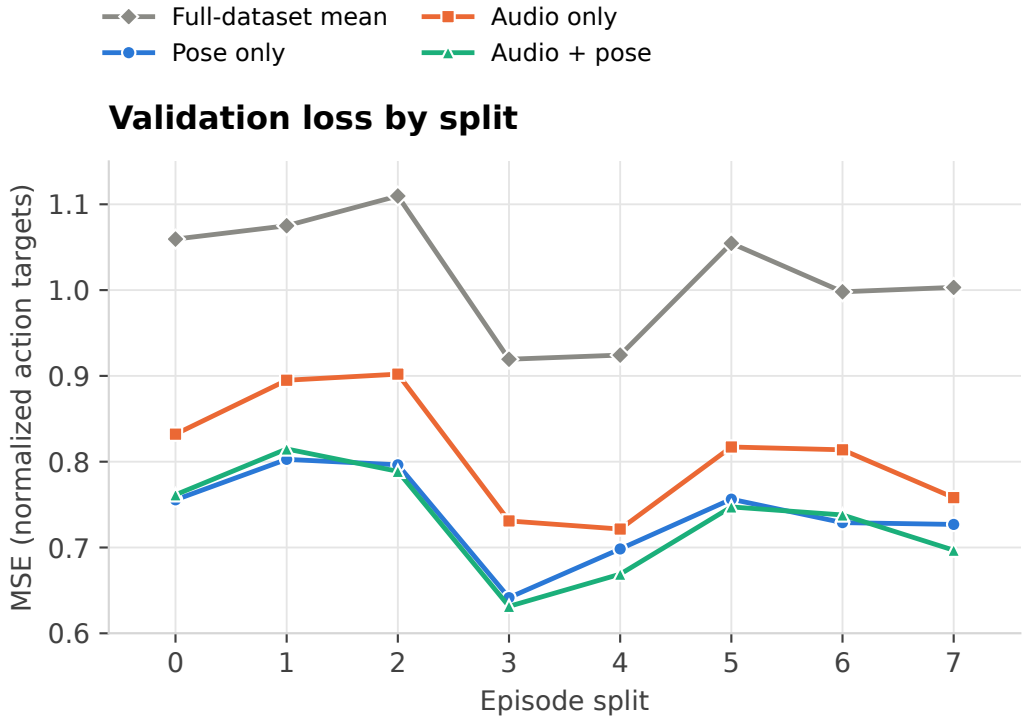
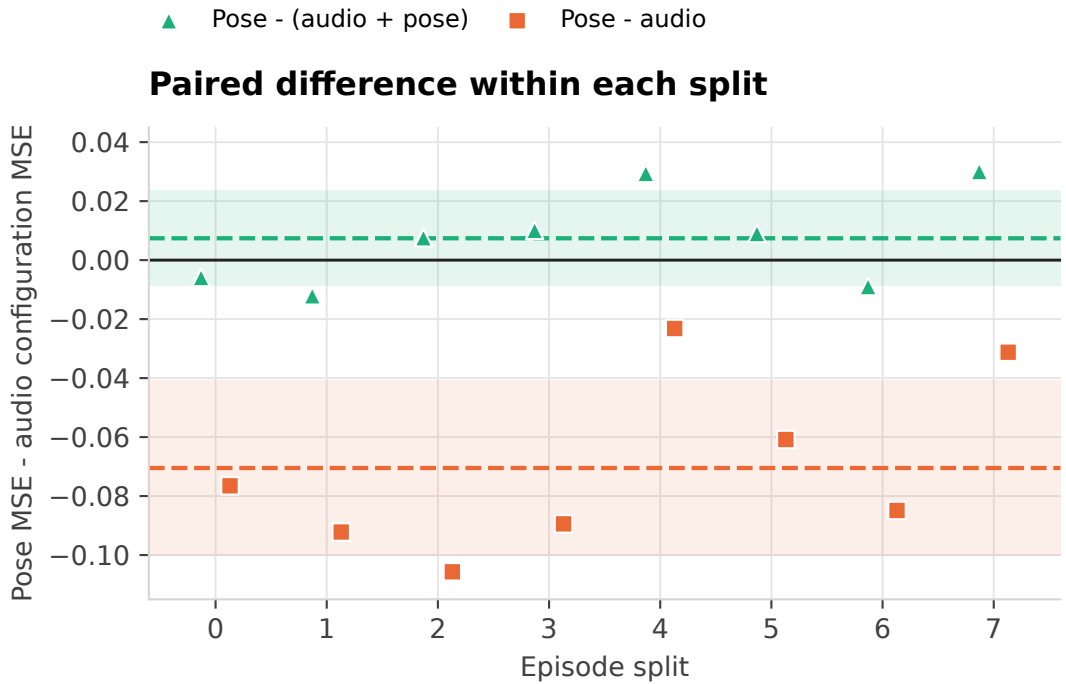
ManiWAV wipe: does audio add to pose?

Frozen AST embeddings | 78 episodes | 16 held out per split





Paired MSE difference: mean +/- sample SD

Pose - (audio + pose)	+0.0074 +/- 0.0162	Audio configuration better in 5/8 splits
Pose - audio	-0.0705 +/- 0.0297	Audio configuration better in 0/8 splits

Positive differences favor the audio configuration. Dashed lines: means; shaded bands: +/-1 SD.

Exploratory scores: best validation epoch of 40; normalization uses all episodes.

Audio alone is worse in 8/8 splits. The added value of audio alongside pose remains inconclusive.